

The Case for Cooperative Resource Sharing for Parallel Applications on Manycore Processors

Paper #175

NN pages

NOTE: please see README.md for list of unincorporated changes

Abstract

The continuous rise of parallel and cloud computing has created new challenges for the systems community. On one hand, high-performance parallelized software requires close collaboration with the operating system and knowledge of underlying hardware. On the other, cloud computing environments attempt to provide an abstraction of a single-tenant system, typically by hiding those same details. This tension will get worse over time as application developers are pressured by the increasing number of cores available in server hardware.

Thus far, the systems community has prioritized maintaining the single-tenancy abstraction at the expense of global application performance. Server virtualization technologies like containers and virtual machines have become popular ways of maintaining isolation between applications that share hardware, but provide little means for cooperative resource sharing between applications.

In this world, developers seeking to maximize performance of their applications have two bad choices: either over-provision units of execution and run the risk of unmanaged resource contention or under-provision them and fall short of performance potential.

We propose Application Resource Schedulers (ARSs), a mechanism for structured communication between the operating system and user-level applications for the purposes of cooperative resource sharing. ARSs allows applications to dynamically alter their behavior in response to changing resource availability, maximizing global and local application performance. We evaluate the feasibility of ARSs by using

one to coordinate CPU usage across parallel execution of applications in the PARSEC benchmark, improving system throughput by up to 60%.

1. Introduction

In recent years, economic and physical realities have pushed chip designers, application developers, and system administrators to rely increasingly on parallelism to satisfy performance demands [22].

Historically, the simplest way to increase application performance was to increase the clock speed of the running CPU. In general, increases in clock speed were driven by shrinking the die on the chip. However, the industry has reached the limits of this strategy: issues such as excess heat dissipation and quantum tunnelling have slowed the pace of miniaturization [16].

As a result, the semiconductor industry has been turning to multiprocessing designs to deliver steady performance gains. For over a decade now, multi-core processors have been widely available, utilizing the increasing number of transistors on a chip without a concomitant increase the CPU clock rates. As an example, Intel introduced its first dual-core processor in 2005 [14]. 9 years later, Intel introduced a multiprocessor with 60 cores capable of running 240 threads in parallel [15].

However, parallel programs struggle to keep pace with the number of cores available, leading to a second trend: an increase in multi-tenancy in order to fully utilize processing hardware.

1.1 Cloud computing and multi-tenancy

Coincident with the the rise of hardware parallelism has been a shift to cloud-based computing. The availability of high-capacity networks, commodity hardware, and virtualization technologies made it attractive for many users to outsource their computing infrastructure to a specialized provider.

To achieve the economies of scale necessary to make cloud computing feasible, service providers commonly share hardware resources among many unrelated applications, while providing the illusion that each application is the sole user of its machine.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

SOSP'17 October 29–31, 2017, Shanghai, China

© 2017 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ISBN [to be supplied]...\$15.00

This illusion is maintained with the help of hardware and software virtualization. Technologies such as virtual machines and containers attempt to hide the multi-tenant nature of the system from users by further abstracting away details of the underlying hardware.

1.2 Needs of high-performance parallelism and cloud computing environment are in tension

Any resource allocation scheme for the cloud requires a delicate balance between application and service provider needs.

Individual applications would like maximum performance for cost. In highly parallel applications, this requires intimate knowledge of the underlying hardware and operating system resources, as well as the guarantee that other actors will not interfere with them.

Cloud providers would like to maximize resource sharing to keep hardware costs low. In addition, they would like high global throughput in order to increase the number of new users they can take on.

These desires are fundamentally at odds. Were each individual application to have its way, they would receive a dedicated server, increasing costs beyond what is feasible for the provider. Were the cloud provider to have their way, application performance would suffer greatly.

This has resulted in an uneasy truce between users and providers. Applications that are sensitive to noisy neighbors over-provision in order to reserve excess capacity, while cloud providers under-provision instances to backstop performance.

1.3 Isolation abstraction limit application performance

Maintaining the illusion of that each application is the sole user of a machine has served operating systems well. However, in a highly parallel, multi-tenant systems, this illusion can inhibit application performance.

Applications designed for high-performance parallelism are typically as "greedy" as possible, attempting to take full advantage of known hardware resources. When many such applications share the same hardware resources, contention for system resources increases, causing unnecessary overhead and causing both individual and total application throughput to suffer.

1.4 The case for cooperative resource sharing

For the most part, research on multi-tenancy in the cloud has borrowed heavily from earlier work on multiprogramming in operating systems [17]. Particular emphasis has been placed on suitable ways to preempt or throttle bad actors in the system to maintain the illusion of isolation [24].

However, modern cloud environments do not have the same characteristics as general-purpose operating systems. Often, servers run many copies of the same application for different users, with the application code being under the full control of the server owner.

Moreover, highly parallel computing workloads often do not benefit from systematic "hiding" of underlying hardware characteristics. As discussed above, parallel programs seek make full use of existing cores, and often assume they have exclusive access to these resources.

In such an environment, there is opportunity for a cooperative resource sharing model to enable "denser" multi-tenancy without sacrificing application performance.

Cooperation over resources requires that the operating system communicate information about resource state to user-level applications, allowing cooperating applications to adjust their run-time behavior accordingly.

This paper presents Application Resource Schedulers as a mechanism to implement cooperative resource sharing for highly parallel applications. The ARS has the following components:

- A **message bus** enabling structured communication between the operating system and user-level applications.
- A **scheduler** that allocates resources between different applications and passes messages to them through the communication channel.
- A **client library** that allowed applications responds to scheduler messages by adjusting runtime behavior.

1.5 Contributions

Based on these observations, we present the following contributions:

1. Observing that current operating system abstractions do not handle contention gracefully when running multiple parallel applications on the same manycore processor and that a cooperative approach may be more suitable
2. Making the case for a structured and dynamic communication mechanism between applications and the operating system to restrict parallel applications and adjust the operating system scheduler
3. Demonstrating that a cooperative approach with the right mechanism can improve system throughput by up to 60%.

2. Application isolation abstractions cause resource contention

This section characterizes the performance of parallel applications under different levels of parallelism, paying particular attention to program behavior when parallelism is high. Based on this characterization, we argue that the existing isolation abstractions are inadequate for multi-tenant, highly parallel workloads.

While many classes of applications might benefit from cooperative resource sharing, this paper focuses its evaluation on highly parallel applications. We chose to focus on parallel applications because:

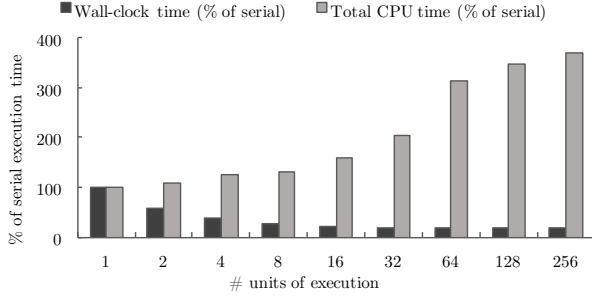


Figure 1: Wall-clock and total CPU execution times of the *bodytrack* application from the PARSEC benchmark suite [2]. Increasing parallelism typically improves wall-clock execution time. However, total CPU time can increase, using more CPU resources to complete the same job. As an example, increasing from 16 to 32 threads increases total CPU time by approximately one-third, but we only see negligible improvement in wall-clock time.

1. Cloud computing workloads are commonly parallel. Due to trends in commodity hardware, parallelizing work will even more important in the future.
2. Parallel workloads uniquely benefit from detailed knowledge of the underlying hardware. They would like to utilize all cores if possible, and often need to micro-optimize to maximize cache hits.

When optimizing parallel programs, application developers must carefully balance the performance benefits of increased concurrency with the performance costs of interference between concurrent activities.

Therefore, deciding how much parallelism to introduce in a program is one of the most important decisions an application developer can make. While it is obvious why under-using parallelism sacrifices performance, the “upper bound” is less understood.

2.1 Performance gains from parallelization are highly variable due to resource contention

In a perfectly parallelizable workload, we would expect to see wall-clock execution speed up linearly with increasing parallelism. In real-world situations, however, wall-clock time typically speeds up sub-linearly with increasing parallelism (see Figure 1).

Clearly, increasing parallelism comes with costs. These costs are well documented, and include:

- **Hardware resource contention.** CPU cores, main memory, bus throughput, etc.
- **Context switch overhead.** Kernel time, cache invalidation, etc.
- **Synchronization overhead.** Common primitives (mutexes, semaphores, etc.) introduce overhead when contended for.

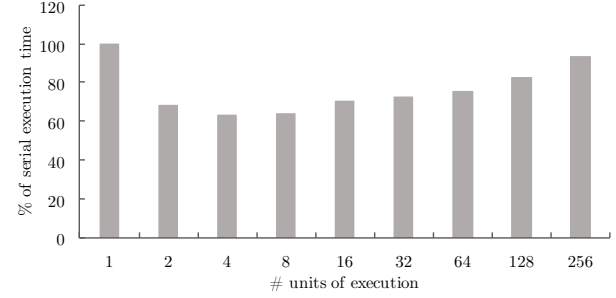


Figure 2: Wall-clock execution times for a job consisting of 16 instances of the PARSEC benchmark application *fluidanimate*. In each observation, we start the 16 instances simultaneously, all with the same number of threads (varied in along the x-axis). Parallel applications are designed to perform better with more units of execution, but performance can quickly degrade under contention.

For most embarrassingly parallel problems, synchronization overhead is not high, as mutex contention is minimal. Context switch overhead can largely be reduced to a resource contention problem, as the cost of a context switch is dominated by cache invalidation [18]. Thus, we focus on the costs stemming from hardware resource contention.

2.2 Isolation abstractions cannot solve resource contention

We argue that the existing abstractions are inadequate for handling resource contention in multi-tenant, highly parallel workloads. The illusion of a single-user machine causes applications to systematically overuse finite hardware resources. As long as applications can expect to be “greedy” with regard to the hardware resource they are presented with, they will behave badly in multi-tenant environments (see Figure 2).

2.2.1 Application doesn’t know what resources it has

In the existing scheme, user-level applications have the power to spawn threads and subprocesses, but little information about the state of hardware resources. Often, the degree of parallelism is statically determined by the application developer, with little regard for portability or the operating environment.

When the number of such applications on a given machine increases, the concurrent applications may not run efficiently. For example, on an Intel Xeon Phi processor with 64 cores, each supporting up to 4 hyperthreads, an application would create 256 threads for this hardware. If its available hardware can only run a small number of threads because several applications are sharing the same machine, the overheads of 256 threads introduce inefficiency.

A common strategy to mitigate this problem to use a sophisticated dynamic parallelism library, such as Intel’s TBB or Microsoft’s TPL [28][11]. However, this only performs

optimization at the application level, running afoul of the scenario described above. Without understanding of the underlying resource environment, any application-local optimization cannot make globally optimal decisions.

2.2.2 How applications are scheduled dramatically affects overall system performance

The second reason is that the kernel does not have enough information to schedule multiple applications to maximize the efficiency or throughput of the hardware. As an example, suppose an application has the following elapse times and speedups:

- 200 seconds on a single core.
- 12.5 seconds on 32 cores (16X speedup), with the best number of threads for 32 cores
- 10 seconds on 64 cores (20X speedup), with the best number of threads for 64 cores.

If the OS runs this application twice on the 64 core, one instance at a time, it will take $10 + 10 = 20$ seconds. However, if it runs two instances concurrently and each instance runs the best number of threads for 32 cores, it will take 12.5 seconds. The throughput of the system improves by 37.5%.

2.3 Current trends in virtualization will exacerbate this issue

Our arguments also apply to virtual machines and containers. Virtual machine monitors provides virtualized hardware for an entire software stack including multiple operating systems. Applications and their operating systems share the same abstractions as applications might on real hardware.

Containers provides isolated environments in a traditional operating systems rather than providing virtualized hardware. Within the context of a container environment, the abstractions between applications and their operating systems remain the same.

3. The case for cooperative resource sharing

Given challenges observed in the previous section, we advocate an extension to the traditional application-OS abstractions with a dynamic, bi-directional communication mechanism between applications and the OS scheduler, as shown in Figure 4.

As long as applications are architected with the assumption that they have exclusive access to machine resources, resource contention is inevitable in parallel, multi-tenant environments. Developers will be required to over-provision compute capacity in order to avoid resource contention, and otherwise program and plan with the expectation that they are sharing a machine with others. They will not be fooled by the illusion that they the sole users of the system.

We argue for cooperative resource sharing as a preferable multi-tasking paradigm. In this scheme, applications are

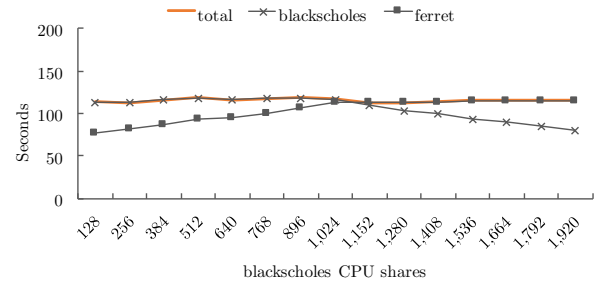


Figure 3: Improving overall performance is difficult without true cooperation from applications. In this simple example, we have a total of 2048 shares of CPU time that we divide among two applications, `blackscholes` and `ferret`, that start at the same time. We use Linux control groups to enforce that the two applications will have roughly the same ratio of time on CPU as the ratio of their weights. We can manipulate an application’s resources with CPU shares, but we cannot alter application behavior via this mechanism and as a result, minimally impact the total execution time (orange, for those with color).

written with the full knowledge that they are co-tenants sharing resources, and agree to allow a centralized scheduler to control their resource consumption. We propose Application Resource Schedulers as a potential mechanism to implement cooperative resource sharing.

3.1 Cooperative resource sharing may uniquely overcome resource contention problems

When applications cooperate, they receive structured information about “real” system resources and instructions about how to utilize resources for greatest effect. In any system where isolation is a core guarantee, this information is by definition unavailable to the application.

Attempts to solve resource contention problems by strengthening isolation guarantees ultimately run up against this limitation. Slicing applications through virtualization does not change fundamental hardware facts (see Figure 3).

3.2 Modern cloud computing environments can limit the downsides of cooperation

The fundamental flaw in cooperative multitasking schemes is that they rely on individual applications to behave themselves. A single bad actor in the system can monopolize resources, which creates a prisoner’s dilemma where it is simpler for the scheduler to pre-empt tasks to ensure fairness.

However, in today’s managed compute environments, all non-trivial applications running on a server are under the full control of the service’s operator. For example, a multi-tenant database-as-a-service like Amazon’s DynamoDB [7] will consist of multiple tenants, all running proprietary Amazon code. In such an environment, ensuring that all applications

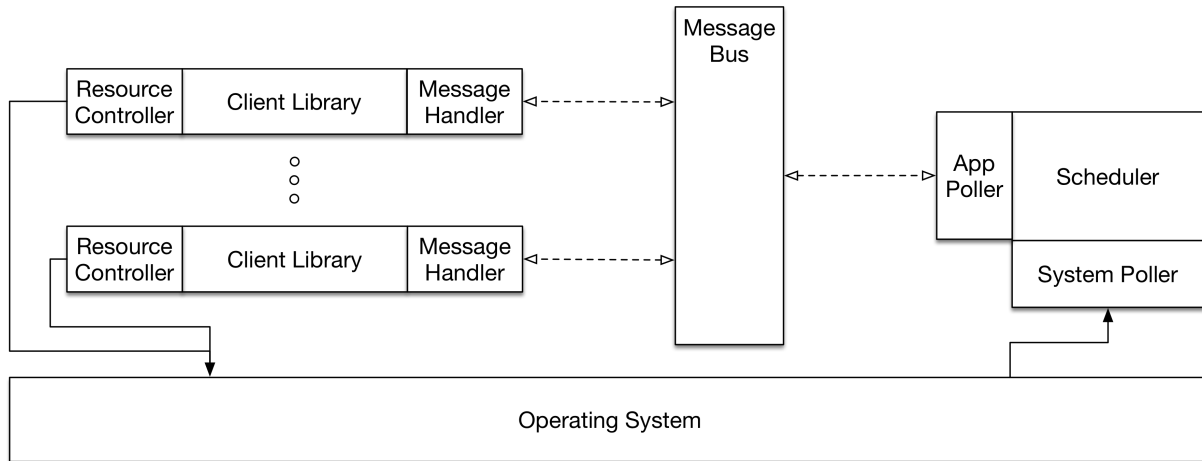


Figure 4: The design for a Application Resource Scheduler, a proposed implementation of cooperative resource sharing

”play by the rules” is simply a matter of integrating some mechanism like an ARS into the application code.

3.3 What are scenarios under which we might use ARSs?

In general, any environment that requires multitasking over several highly parallel applications could benefit from cooperative resource sharing. Some examples include:

- **OLAP and batch processing workloads.** For big data analytics and data warehousing applications, many highly parallel jobs may be placed a single node. The application itself can be modified to cooperate with the ARS, and the interest of the service owner is likely to maximize global throughput per node.
- **Proprietary compute clouds.** Many companies operate private clouds, where all users and applications are under

4. Designing a CPU Application Resource Scheduler for parallel programs

This section discusses the design of a Application Resource Scheduler, a system to enable cooperative resource sharing. This design is generalizable to any type of shared resource and any type of application, but in the interest of brevity we will discuss it in terms of sharing CPU resources between highly parallel applications.

4.1 Message bus

This component enables communication between the scheduler and client libraries hosted on each cooperating application.

It is important to note that, as a single component accessible to all applications, the message bus is itself a shared resource. Therefore, care must be taken that it does not become a bottleneck in the system.

To mitigate contention for this component, the scheduler and client libraries implement a stateful, asynchronous pro-

tolocol over this message bus that allows them to operate independently of one another, decreasing the likelihood of synchronization overhead.

4.2 Scheduler

The scheduler is responsible for maintaining global resource state and allocating resource between cooperating applications in accordance with some predetermined policy.

In our example of a CPU-sharing ARS, the scheduler will decide how many software threads each application should be allowed.

4.2.1 System poller

To collect and store this information, the scheduler maintains a system poller that communicates with the operating system about hardware resources and stores the resulting data in a time-series data structure.

In the case of a CPU-sharing ARS, global resource state could, for example, consist of:

- Number of cores and degree of hyperthreading
- Number of runnable processes
- Number of sleeping processes

4.2.2 Application poller

At the same time, the scheduler continuously checks for messages on the message bus. These messages provide confirmation that cooperating applications have indeed adjusted their behavior in accordance with the scheduler, as well as passing along any application-level data that the scheduling policy might require.

Application-level data is likely to change over the lifetime of application. This is especially true of applications with distinct phases of operation. Thus, it is sensible to store application-level data in a time-series data structure as well.

In the CPU-sharing case, the application would send the number of software threads it is operating with. Additional

data shared could be, for example, application-specific performance instrumentation or user-specific SLAs.

4.2.3 Scheduling policy

Based on data gathered from the system poller and the application poller, the scheduler makes a decision on resource allocation to cooperating applications based on its scheduling policy. These allocations are then communicated via the message bus to each application's client library.

The efficiency information of each application can be used by the policy to achieve certain goals such as maximizing throughput, minimizing average latency, or maximizing hardware utilization.

There is a wide body of operating systems research and industry work on the merits of various scheduling policies [19]. We note that they could be implemented in any ARS scheduler.

4.3 Client library

The client library is hosted by all cooperating applications. In a typical cloud computing environment, all applications of non-trivial resource consumption will be cooperating, since the environment is under the full control of the system operator. In the case of as CPU-sharing ARS, the client library has two modules:

4.3.1 Message handler

This module periodically checks the message bus for any messages from the scheduler about how to adjust the application's resource consumption, and updates the client library's internal state accordingly.

In the CPU-sharing case, the message from the scheduler takes the form of a recommendation for how many software threads to use.

4.3.2 Resource controller

This module takes the scheduler's recommendation about resource consumption and executes it by adjusting run-time application behavior. In the CPU-sharing case, this would take the form of a thread pool which supports dynamic size adjustment.

This is the module that is most tightly coupled with application logic. As a result, care must be taken to design the controller in a way that is easily integrated into existing systems. For example, the CPU-sharing resource controller could wrap a popular existing thread pool implementation, allowing clients to specify a starting size with the understanding that it may change over time. As another example, a memory-sharing resource controller could implement a dynamically-sized arena that applications use in place of ad-hoc memory allocation.

5. Implementation considerations

In this section, we discuss an example of a concrete implementation of a CPU-sharing Application Resource Sched-

uler design. We begin by outlining some of our general design goals, and proceed by examining implementation considerations for each component.

5.1 Design goals

At a high level, we wanted our implementation to conform to the following goals:

- **Low overhead.** As the primary purpose of cooperative resource sharing is to increase global application performance, care must be taken that our ARS implementation not become a practical bottleneck.
- **Graceful fallback.** To ease application integration, simplify testing, and increase robustness, our implementation should fall back to sensible default in the event that any component becomes unavailable.

In addition, we had the following non-goals:

- **Exhaustive performance optimization.** Our goal was to make the case that cooperative resource sharing can improve performance for highly parallel workloads, not to tune a production environment for maximum efficiency.
- **Sophisticated scheduling policy.** Significant work has already been done in scheduling theory and practice. Our implementation need not retread existing ground to prove that cooperation is desirable.
- **Intrusive resource control.** While it may be tempting to specify elaborate resource acquisition and consumption requirements for applications, in practice this would require that developers substantially revise their computation models in order to integrate with our library.

5.2 Message bus

As a globally shared resource, message bus must be fast even in the presence of many clients. In practice, this means any implementation must achieve:

- **Fast reads and writes.** An interprocess communication scheme must be chosen such that both the scheduler and client libraries have extremely fast access.
- **Low synchronization overhead.** Even with many (hundreds) of applications, there should be little or no contention for synchronization resources..

To achieve fast reads and writes, we implement the message bus using shared memory as our communication medium. We create two channels in the allocated shared memory, corresponding to each direction of communication.

The scheduler-to-application channel contains messages regarding the scheduler's latest recommendation for the number of software threads for applications to use. The application-to-scheduler channel contains messages where applications register their existence with the scheduler and update their current thread count.

Both channels are protected by reader-writer locks to ensure that multiple readers do not contend for access to the channel. In our implementation, the client library only sends and receives messages every 500 milliseconds, which makes lock contention negligible in practice. If the latency of client response is a concern, finer-grained synchronization models are feasible (for example, giving each application its own channel).

5.3 Scheduler

5.3.1 Scheduling policy

As sophisticated scheduling was a non-goal, we intentionally made our scheduler policy as minimal as possible. Our policy only requires one piece of system information to make allocation decisions: the number of runnable processes currently waiting to run. On this basis, a single recommendation is issued for all applications to follow. No distinctions are made between applications, and the scheduling policy only tries to maximize global application throughput.

5.3.2 System poller

For low overhead communication between the kernel and userspace, we rely on `procfs`. This allows the system poller to gather basic information about the system without the overhead of a system call or some dynamic tracing method. Our implementation only reads the `procs.running` field of `/proc/stat`, which gives the number of threads currently running on CPU.

This information is gathered once per second and saved in a fixed-size time series data structure.

5.3.3 Application poller

Our simple scheduling policy does not require information from the application. While we implemented a channel for the application to pass messages to the scheduler about the current number of threads it is using, this information is unused by the scheduler.

5.4 Client library

In implementing the client library, a decision needs to be made about when the resource controller will intervene to set the number of threads in the application. There are three options, each with varying levels of intrusion into the application logic:

1. **Static parallelism.** This is equivalent to no resource controller at all; the number of threads is determined by user input.
2. **Dynamic parallelism.** On start-up, the client library consults the scheduler for how many threads to use. This number is kept throughout the lifetime of the application.
3. **Dynamic run-time parallelism.** Throughout the lifetime of the application, the client library continuously receives messages from the scheduler and resizes its thread pool accordingly.

While dynamic run-time parallelism is the optimal resource control model from a performance perspective, it is highly intrusive and often requires the application developer to re-architect the concurrency model of their program.

For example, the `canneal` benchmark in PARSEC uses a highly aggressive lock-free synchronization strategy that recovers from data races rather than try to avoid them [2]. Changing the number of threads dynamically would require substantial effort, and repartitioning the work when changing thread count would require global locking that would sacrifice the performance benefits of a lock-free strategy.

Thus, while we implemented the features necessary for dynamic run-time parallelism in our client library, we opted to include a mode that disabled the run-time thread count updates.

6. Performance implications

In our evaluations we seek to answer whether cooperative scheduling can meaningfully improve overall performance. Unless otherwise noted, we focus on overall system throughput. The workloads that might benefit most from our proposal—batch data analytics—motivate this choice.

6.1 Prototype implementation

When implementing our prototype, we opted for simpler designs and heuristics. The message bus, scheduler, and client library together comprise fewer than 2000 lines of C and C++. Thanks to the design of the PARSEC benchmark suite, we were able to integrate the client library into many of the benchmark applications with fewer than 100 lines of modified code. To simplify communication among components, we use shared memory and locks, allowing us to forgo message queues or additional threads to maintain communication.

Although we implemented more complex mechanisms, we were able to demonstrate strong results using the most basic configuration. In the following experiments, the scheduler only polls the system-wide run queue and ignores any other system or application-specific data. Given that the target workloads are CPU-intensive, using thread contention as a proxy for resource contention is reasonable. Despite this simplicity, however, the application implicitly gives the scheduler a lot of information by registering as a client. Most importantly, the application signals that it is capable of dynamic run-time parallelism and is willing to follow directives from the scheduler.

During run-time, the scheduler provides applications with a recommended number of threads to create on start. We assume that applications are well-behaved, but it would not be difficult to implement punitive measures.

6.2 Experimental environment

We conducted experiments on four different machine types: a quad-core desktop (8 hardware threads), a dual-socket 9-

Name	Application domain	Parallelization		Working set	Data usage	
		Model	Granularity		Sharing	Exchange
*blackscholes	Financial analysis	data-parallel	coarse	small	low	low
bodytrack	Computer vision	data-parallel	medium	medium	high	medium
*canneal	Engineering	unstructured	fine	unbounded	high	high
*dedup	Enterprise storage	pipeline	medium	unbounded	high	high
facesim	Animation	data-parallel	coarse	large	low	medium
ferret	Similarity search	pipeline	medium	unbounded	high	high
*fluidanimate	Animation	data-parallel	fine	large	low	medium
freqmine	Data mining	data-parallel	medium	unbounded	high	medium
*streamcluster	Data mining	data-parallel	medium	medium	low	medium
*swaptions	Financial analysis	data-parallel	coarse	medium	low	low
vips	Media processing	data-parallel	coarse	medium	low	medium
x264	Media processing	pipeline	coarse	medium	high	high

Table 1: Applications annotated with a * have been integrated with the ARS client library. This is a qualitative summary of all the workloads in the PARSEC benchmark. PARSEC workloads were chosen to cover different application domains, parallel models and runtime behaviors. [3]

core c4.8xlarge Amazon EC2 instance (36 vCPUs), a 4-socket 10-core workstation (80 hardware threads), and a cluster with many 68-core machines (272 hardware threads). Although our findings are consistent across machines, we choose to present results from the 68-core machines to illustrate the challenges introduced by these manycore chips.

The cluster is part of a research installation generously maintained by the Intel® Corporation and features 256 nodes each equipped with an Intel Xeon Phi™ CPU (Model 7250). Each Xeon Phi™ has 68 physical 1.4GHz cores with 4 hardware threads per core. Each physical core has 32K of L1d and L1i cache and shares a 1024K L2 cache with one other core. The CPU can be configured in several ways. First, we use the fully symmetric NUMA topology as opposed to the quadrant mode to simplify the profiling and debugging process. Second, we use the 16GB on-package memory as a large last-level cache rather than use it as ultra-fast memory so our results are more generally applicable to less sophisticated hardware.

6.3 Benchmark applications

To evaluate our ARS, we use a popular parallel application suite, PARSEC [2], to generate a large number of diverse workloads spanning several classes of parallel programs and problem domains (see Table 1). After integrating the client library into each application, we control their levels of parallelism as they enter the system. For six applications (see starred applications in Tabl 1), we also instrumented the capability to dynamically adjust parallelism during program execution. However, in experiments, we found that the overhead of such a design is only justified for longer-running applications residing with unpredictably busy neighbors.

6.4 Baseline comparisons

In order to create an apples-to-apples comparison unencumbered by differences in implementation, we chose to evaluate the workloads using two existing thread libraries with and without cooperation. We use the Pthreads standard as an example of lower-level parallel programming and the Intel® Thread Building Blocks (TBB) library as an example of a sophisticated industrial-grade library. Eleven of the twelve PARSEC benchmark applications are implemented using Pthreads and five also have TBB implementations.

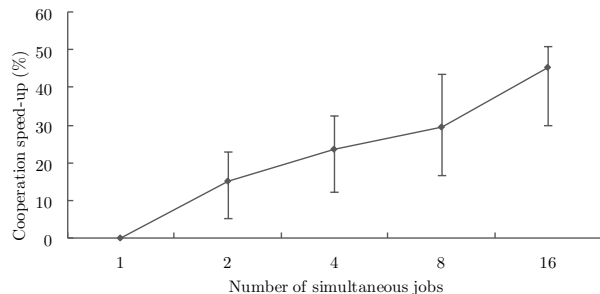
6.5 Generating workloads

In order to compare performance, we create several hundred randomly-generated workloads by drawing samples from the set of benchmark applications. We focus our attention on constant load scenarios as these are most advantageous to the existing parallel programming libraries. More dynamic loads are amenable to our approach since they offer many opportunities to change an application’s parallelism to adjust to new circumstances.

We run each workload with Pthreads and TBB in their normal use case and as part of a cooperative environment. For each of these four cases, we run for several levels of cotenancy, modeled by some constant number of simultaneously running applications.

6.6 Coordinating with ARSs

We show our main results in Figure 5. The pair of figures show the same experiments for Pthreads and TBB. The vertical axis indicates how much faster the cooperative approach completed workloads relative to the pre-emptive approach on the same workload. The line shows mean speed-up for the specified load level and the error bars show the maximum and minimum observed speed-ups.



(a) Pthreads implementation



(b) Intel® Threading Building Blocks (TBB) implementation

Figure 5: The benefit from cooperation increases with contention. Each observation is 100 jobs (100 tasks each) run both with and without cooperation. The line represents the mean throughput (time to job completion) while error bars represent the maximum and minimum speed-ups we observed. Of note, we did not observe our mechanism harming performance.

There are two important observations. First, we did not observe in any scenario where cooperating led to a worse outcome. Second, cooperation improves performance to a greater extent with increasing contention. We were able to achieve such a result with a simple design because our workload has a number of structural features (e.g., CPU-intensive, parallel structure) that lend themselves to cooperation.

7. Discussion and related work

There are three broad areas of related work. First, there is a long and significant history of coordinating resources among applications. Second, there are several parallel programming libraries that touch on many similar ideas (e.g., monitoring resources, manipulating application behavior). Finally, from a broader perspective, there are the works that investigate and tease the application-OS interface.

7.1 Coordinating resource consumption

Coordinating resources is an important problem in many domains. In distributed systems alone, we have seen numerous takes on the problem [31][1][27][6][30][13][8][12]. We acknowledge that we could not possibly mention all relevant approaches. Instead, we select a few diverse examples that specifically address a single node. We classify these approaches approximately as application-aware and hardware-aware.

Application-aware approaches have typically relied on application directives (e.g., `UNLIKELY`, processor affinity, DSLs like OpenCL), online monitoring, offline profiling, or some combination of the three. Gang scheduling [25] marks one of the earlier approaches to coordination, describing a mechanism for communicating units of execution to align and avoid blocking. Cache- [26] and contention-aware [4] schedulers examine historical application behavior to make intelligent resource allocations. Finally, the rich literature of real-time systems [20][9] demand very precise information from applications.

However, we observe that while these approaches make use of application information, they make decisions without the cooperation of applications. We believe that cooperative multitasking is most similar in spirit to our approach, but as evidenced by virtually all operating systems, it may have chosen dangerous abstractions.

By contrast, hardware-aware approaches attempt to probe information about underlying hardware [5] [29] to make resource decisions. Perhaps most famously the Linux scheduler [23] shifted from application-aware, with its complex system for classifying applications as interactive or batch, to hardware-aware (e.g., via scheduling domains). While the Linux scheduler is remarkably effective at its unenviable task as a generic scheduler, it has its own challenges[21]. We argue that the Linux scheduler need not take on additional complexity and that we can instead push responsibility to the applications themselves.

7.2 Parallel programming libraries

Many parallel programming libraries exist to ease the burden of concurrency on developers. While some might be considered more syntactic sugar, many are robust software projects in their own right, handling thread creation, destruction, scheduling, and more. However, to the authors' knowledge, we present the first POSIX scheduler that dynamically scales the level of application parallelism at run-time in response to fluctuating resources. Given our discussion of Intel's TBB, we highlight Microsoft's .NET Thread Pool as another advanced threading library.

The key insight from Microsoft's Common Language Runtime (CLR) .NET thread pool [11][10] is to borrow ideas from control theory rather than observe system metrics that may not be related to application behavior. They monitor throughput of work queues and use closed-loop control and a hill-climbing algorithm to determine whether to add or remove threads. However, one of the primary motivations for this mechanism was the expense of maintaining unused threads in Windows.

8. Conclusion

8.1 Summary

We have argued the case for cooperative resource sharing, and presented an example implementation of it in the form of Application Resource Schedulers. In a world where parallelism and multi-tenancy are both increasing, the traditional abstraction that each application is the sole user of the system is becoming untenable. Attempts to solve this problem have hitherto tried to provide even stronger isolation guarantees with limited success under heavily parallel workloads.

We argue that weakening the isolation abstraction allows for applications to cooperate over the utilization of shared resources. By evaluating an implementation of cooperative resource sharing (ARSSs) against the PARSEC parallel computing benchmark, we have shown that even naive cooperation over resources can yield performance benefits.

8.2 Limitations

This paper primarily seeks to make a case for a new direction in research, not to provide exhaustive proof that cooperative resource sharing is the optimal paradigm. To that end, our implementation is simple, and does not boast many basic features one would want in a production-ready system.

Furthermore, the traditional downsides of cooperative multitasking models still apply. Malicious applications (or even just uncooperative ones) will likely ruin system performance. While these downsides are limited in many multi-tenant environments, there are many domains in which cooperation is not appropriate.

Lastly, we have only examined a specific kind of workload in a highly multi-tenant environment. The PARSEC benchmark aims to be a diverse representation of CPU-bound parallel workloads, but it is possible that our implementation over-fits to a synthetic benchmark and may not perform as well in a "real world" environment.

8.3 Future work

As cooperative resource sharing is relatively unexplored, substantial work needs to be done to validate that it is feasible under different workloads and with different resources under contention. More sophisticated scheduling policies can be tested to see whether cooperation can make the same kinds of tradeoffs about fairness, latency, and so on as a more traditional scheduling approach. Dynamic run-time parallelism should be evaluated against simpler approaches to application resource control, and concurrent programming paradigms that better support dynamic run-time parallelism should be demonstrated.

References

- [1] BHATTACHARYA, A. A., CULLER, D., FRIEDMAN, E., GHODSI, A., SHENKER, S., AND STOICA, I. Hierarchical scheduling for diverse datacenter workloads. In *Proceedings of the 4th annual Symposium on Cloud Computing* (2013), ACM, p. 4.
- [2] BIENIA, C., KUMAR, S., SINGH, J. P., AND LI, K. The parsec benchmark suite: Characterization and architectural implications. In *Proceedings of the 17th international conference on Parallel architectures and compilation techniques* (2008), ACM, pp. 72–81.
- [3] BIENIA, C., AND LI, K. Characteristics of workloads using the pipeline programming model. In *Computer Architecture* (2012), Springer, pp. 161–171.
- [4] BLAGODUROV, S., ZHURAVLEV, S., AND FEDOROVA, A. Contention-aware scheduling on multicore systems. *ACM Transactions on Computer Systems (TOCS)* 28, 4 (2010), 8.
- [5] BLAGODUROV, S., ZHURAVLEV, S., FEDOROVA, A., AND KAMALI, A. A case for numa-aware contention management on multicore systems. In *Proceedings of the 19th international conference on Parallel architectures and compilation techniques* (2010), ACM, pp. 557–558.
- [6] CHOWDHURY, M., ZHONG, Y., AND STOICA, I. Efficient coflow scheduling with varys. In *ACM SIGCOMM Computer Communication Review* (2014), vol. 44, ACM, pp. 443–454.
- [7] DECANDIA, G., HASTORUN, D., JAMPANI, M., KAKULAPATI, G., LAKSHMAN, A., PILCHIN, A., SIVASUBRAMANIAN, S., VOSSHALL, P., AND VOGELS, W. Dynamo: amazon's highly available key-value store. *ACM SIGOPS operating systems review* 41, 6 (2007), 205–220.
- [8] GHODSI, A., ZAHARIA, M., HINDMAN, B., KONWINSKI, A., SHENKER, S., AND STOICA, I. Dominant resource fairness: Fair allocation of multiple resource types.
- [9] GOYAL, P., GUO, X., AND VIN, H. M. A hierarchical cpu scheduler for multimedia operating systems. In *OSDI* (1996), vol. 96, pp. 107–121.
- [10] HELLERSTEIN, J. L. Configuring resource managers using model fuzzing: A case study of the .net thread pool. In *Integrated Network Management, 2009. IM'09. IFIP/IEEE International Symposium on* (2009), IEEE, pp. 1–8.
- [11] HELLERSTEIN, J. L., MORRISON, V., AND EILEBRECHT, E. Optimizing concurrency levels in the .net threadpool: A case study of controller design and implementation. *Feedback Control Implementation and Design in Computing Systems and Networks* (2008).
- [12] HENDERSON, R. L. Job scheduling under the portable batch system. In *Workshop on Job Scheduling Strategies for Parallel Processing* (1995), Springer, pp. 279–294.
- [13] HINDMAN, B., KONWINSKI, A., ZAHARIA, M., GHODSI, A., JOSEPH, A. D., KATZ, R. H., SHENKER, S., AND STOICA, I. Mesos: A platform for fine-grained resource sharing in the data center. In *NSDI* (2011), vol. 11, pp. 22–22.
- [14] INTEL CORPORATION. *Intel's first dual-core processor-based platform*, 2005.
- [15] INTEL CORPORATION. *Intel Xeon Phi Coprocessor 3120A*, 2013.
- [16] KISH, L. B. End of moore's law: thermal (noise) death of integration in micro and nano electronics. *Physics Letters A* 305, 3 (2002), 144–149.

- [17] KREBS, R. *Performance Isolation in Multi-Tenant Applications*. PhD thesis, Karlsruhe Institute of Technology, 2015.
- [18] LI, C., DING, C., AND SHEN, K. Quantifying the cost of context switch. In *Proceedings of the 2007 workshop on Experimental computer science (2007)*, ACM, p. 2.
- [19] LI, T., BAUMBERGER, D., AND HAHN, S. Efficient and scalable multiprocessor fair scheduling using distributed weighted round-robin. ACM.
- [20] LIU, C. L., AND LAYLAND, J. W. Scheduling algorithms for multiprogramming in a hard-real-time environment. *Journal of the ACM (JACM)* 20, 1 (1973), 46–61.
- [21] LOZI, J.-P., LEPERS, B., FUNSTON, J., GAUD, F., QUÉMA, V., AND FEDOROVA, A. The linux scheduler: a decade of wasted cores. In *Proceedings of the Eleventh European Conference on Computer Systems (2016)*, ACM, p. 1.
- [22] MACK, C. A. Fifty years of moore’s law. *IEEE Transactions on semiconductor manufacturing* 24, 2 (2011), 202–207.
- [23] MOLNAR, I. Cfs scheduler. In *Linux (2007)*, vol. 2, p. 36.
- [24] MULTITENANCY, A. Enforcing multitenancy for cloud computing environments.
- [25] OUSTERHOUT, J. K., ET AL. Scheduling techniques for concurrent systems. In *ICDCS (1982)*, vol. 82, pp. 22–30.
- [26] PHILBIN, J., EDLER, J., ANSHUS, O. J., DOUGLAS, C. C., AND LI, K. Thread scheduling for cache locality. In *ACM SIGOPS Operating Systems Review* (1996), vol. 30, ACM, pp. 60–71.
- [27] REGEHR, J. D. *Using hierarchical scheduling to support soft real-time applications in general-purpose operating systems*. PhD thesis, University of Virginia, 2001.
- [28] REINDERS, J. *Intel threading building blocks: outfitting C++ for multi-core processor parallelism*. ” O’Reilly Media, Inc.”, 2007.
- [29] TOPCUOGLU, H., HARIRI, S., AND WU, M.-Y. Performance-effective and low-complexity task scheduling for heterogeneous computing. *IEEE transactions on parallel and distributed systems* 13, 3 (2002), 260–274.
- [30] WATERMAN, J. *Coordinated Resource Management in Networked Embedded Systems*. PhD thesis, 2012.
- [31] YOO, A. B., JETTE, M. A., AND GRONDONA, M. Slurm: Simple linux utility for resource management. In *Workshop on Job Scheduling Strategies for Parallel Processing (2003)*, Springer, pp. 44–60.